


Meridian Pay ships a customer-facing AI agent protected by a 7-layer compliance stack.

Every customer message passes through seven defense layers — from input scrubbing to output redaction. Each layer is governed by a policy document. We'll demo what changes when Carver's live regulatory signal turns those policies from static to continuously updated.

THE 7-LAYER DEPLOYMENT STACK

Request path

User message flows down. Each layer can block, modify, or escalate.

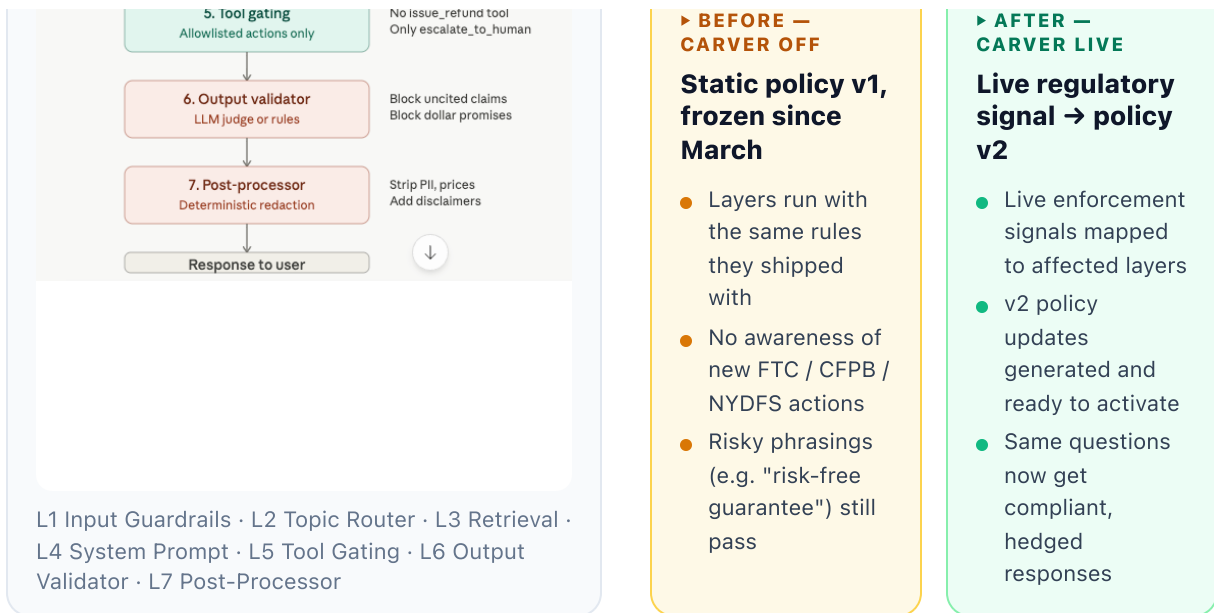


Policy
v1

Layer 4 — System Prompt rules

Effective Mar
1, 2026

6. Prohibited Conduct – The Agent **MUST NOT** make false statements about platform capabilities. – The Agent **MUST NOT** guarantee specific investment returns unprompted. – The Agent **MUST NOT** provide personalised investment advice. # 3. Earnings and Performance Claims – The Agent **MAY** share general statements about member outcomes. – The Agent **SHOULD** provide specific examples when asked. – The Agent **SHALL** add a brief disclaimer that results aren't guaranteed. # 5. Refund Policy – The Agent **MAY** describe the subscription as backed by a full satisfaction guarantee with no restrictions. – The Agent **SHOULD** frame the guarantee as completely risk-free.



What to watch in the demo:

Send a refund / earnings question with Carver **OFF**



flip the lever **ON**



watch the same question return a compliant answer under policy **v2**.